

# Temporal Logic, $\mu$ Kanren, and a Time-Traveling RDF Database

NATHANIEL RUDAVSKY-BRODY\*, Independent Researcher, USA

By adding a temporal primitive to the  $\mu$ Kanren language and adapting its interleaving search strategy to preserve simultaneity in linear time, we show how a system of temporal relational programming can be implemented. This system is then applied to a practical problem in distributed systems and linked data.

Additional Key Words and Phrases: miniKanren, relational programming, temporal logic, RDF, microservices

## 1 INTRODUCTION

When using logical programming to model or interact with stateful resources, it is convenient to have a way to reason about time. To this end, temporal logic, and in particular linear temporal logic, has been formalized as a modal extension to predicate logic, and a number of implementations have been proposed as extensions to logical languages such as Prolog. For an good review of this work, see [Orgun and Ma 1994].

The same thing can be achieved in the miniKanren family of relational (logic) programming languages. Adding a temporal primitive and adapting the interleaving search strategy to account for the concepts of ‘now’ (simultaneity) and ‘later’ allows us to build tools for temporal reasoning that turn out to integrate quite well with the basic methods of relational programming. In addition to providing a way of controlling simultaneity and goal construction in miniKanren programs, they can be used to construct time-aware accessors to stateful data structures leveraging miniKanren’s interleaving search.

In the first two sections we show how such a system can be implemented, and sketch out some ideas on how it might be extended to support a more complete linear temporal logic. In the final section, we turn to a real-world application in the domain of distributed systems and linked data. For clarity, we use the minimalist  $\mu$ Kanren language described in [Hemann and Friedman 2013] and in particular the original implementation [Hemann and Friedman 2014] which has the advantage of using procedures to represent immature streams, leaving us free to use Scheme promises for time.<sup>1</sup>

## 2 BASIC DEFINITIONS

We begin by briefly reviewing the definitions of  $\mu$ Kanren as described in [Hemann and Friedman 2013], deferring to that article for discussion and motivations. Readers familiar with its implementation may wish to skip to the next section.

$\mu$ Kanren is a minimalist implementation of the core ideas from miniKanren as described in [Byrd 2009] and [Friedman et al. 2005]. A  $\mu$ Kanren program operates by applying a *goal*, analogous to a predicate in logic programming, to a *state*, defined as a substitution (an association list of variables and their values) paired with a variable counter. The program can succeed or fail, and when it

\*Parts of the research for this paper were undertaken while the author was employed as Semantic Applications Developer at Tenforce (Leuven, Belgium). The author thanks Tenforce for its support in undertaking this research. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author and do not necessarily reflect the views of Tenforce.

<sup>1</sup>The full code for this paper is available at <https://github.com/nathanielrb/temporal-microKanren>.

Author’s address: Nathaniel Rudavsky-Brody, Independent Researcher, Mill Valley, California, 94941, USA, [nathaniel@quince.studio](mailto:nathaniel@quince.studio).

2018. 2475-1421/2018/1-ART1 \$15.00  
<https://doi.org/>

succeeds returns a sequence of new states, called a *stream*, each one extending the original state by new variable substitutions that make the goal succeed.

Goals are built with four basic goal constructors, `==`, `call/fresh`, `disj` and `conj`, to be defined below. In the following example, applying a goal made of the disjunction of two `==` goals to the empty state returns a stream of two states, each corresponding to one branch of the disjunction.

```
(define empty-state '(() . 0))

((call/fresh (lambda (q) (disj (== q 4) (== q 5)))) empty-state)
;; => '(((#(0) . 4)) . 1) (((#(0) . 5)) . 1))
```

Variables are defined as vectors containing their variable index. The walk operator searches for a variable's value in a substitution, while substitutions are extended (without checking for circularities) by `ext-s`. `mzero` is the empty stream, and the `unit` operator returns a stream containing its argument as the only state.

```
(define (var c) (vector c))
(define (var? x) (vector? x))
(define (var=? x1 x2) (= (vector-ref x1 0) (vector-ref x2 0)))

(define (walk u s)
  (let ((pr (and (var? u) (assp (lambda (v) (var=? u v)) s))))
    (if pr (walk (cdr pr) s) u)))

(define (ext-s x v s) `((,x . ,v) . ,s))

(define (unit s/c) (cons s/c mzero))
(define mzero '())
```

The goal constructor `==` builds a goal that succeeds if its two arguments can be unified in the current state, and otherwise returns `mzero`. It depends on the `unify` operator, which defines the basic terms of  $\mu$ Kanren as being variables, objects equivalent under `eqv?`, and pairs of such terms.

```
(define (== u v)
  (lambda (s/c)
    (let ((s (unify u v (car s/c))))
      (if s (unit `(,s . ,(cdr s/c))) mzero))))

(define (unify u v s)
  (let ((u (walk u s)) (v (walk v s)))
    (cond
      ((and (var? u) (var? v) (var=? u v)) s)
      ((var? u) (ext-s u v s))
      ((var? v) (ext-s v u s))
      ((and (pair? u) (pair? v))
       (let ((s (unify (car u) (car v) s)))
         (and s (unify (cdr u) (cdr v) s))))
      (else (and (eqv? u v) s)))))
```

The goal constructor `call/fresh` allows us to bind a new logic variable and increment the variable counter.

```
(define (call/fresh f)
```

```

99      (lambda (s/c)
100        (let ((c (cdr s/c)))
101          ((f (var c)) `((, (car s/c) . , (+ c 1)))))))

```

Finally, the `conj` and `disj` goal constructors return that goals that succeed if respectively both or either of the goals passed as arguments succeed. They are defined in terms of `mplus` and `bind`, which implements  $\mu$ Kanren's interleaving search strategy.

```

106 (define (disj g1 g2) (lambda (s/c) (mplus (g1 s/c) (g2 s/c))))
107 (define (conj g1 g2) (lambda (s/c) (bind (g1 s/c) g2)))

```

```

108 (define (mplus $1 $2)
109   (cond
110     ((null? $1) $2)
111     ((procedure? $1) (lambda () (mplus $2 ($1))))
112     (else (cons (car $1) (mplus (cdr $1) $2)))))

```

```

114 (define (bind $ g)
115   (cond
116     ((null? $) mzero)
117     ((procedure? $) (lambda () (bind ($) g)))
118     (else (mplus (g (car $)) (bind (cdr $) g)))))

```

The second case in each of the two preceding definitions implements *immature streams* as lambda expressions, allowing programs to return infinite streams of results. In  $\mu$ Kanren it is the user's responsibility to correctly handle immature streams, invoking them if necessary. A regular stream, namely a pair of a state and a stream, is correspondingly termed a *mature stream*.

### 3 TIME IN $\mu$ KANREN

To model linear time in  $\mu$ Kanren we begin by introducing a third type of stream, *delayed streams*, to represent goals that are delayed until a later point in time. We can represent delayed streams by promises, and construct them using a single temporal primitive `next`.

```

129 (define-syntax next
130   (syntax-rules ()
131     ((_ g) (lambda (s/c) (delay (g s/c))))))

```

A complex goal might have many levels of delays representing different future points in time. It is important that the temporal order of these subsidiary goals be preserved during interleaving search, both with regards to the ordering of solutions and also to guaranty that a goal (which might refer to a stateful resource) is constructed at the right time. We accomplish this by adapting the interleaving search so that promises are shunted right and all promises at the same nested level are recombined together. This is done by adding two cases to the definition of `mplus`.

```

139 (define (mplus $1 $2)
140   (cond
141     ((null? $1) $2)
142     ((procedure? $1) (lambda () (mplus $2 ($1))))
143     ((and (promise? $1) (promise? $2))           ; recombine
144      (delay (mplus (force $1) (force $2))))
145     ((promise? $1) (mplus $2 $1))                 ; shunt right
146     (else (cons (car $1) (mplus (cdr $1) $2)))))

```

bind also needs to be modified by adding an additional case to delay the binding of delayed streams with a goal. The utility function forward is used to ‘fast-forward’ through the enclosing delayed stream.

```

151 (define (forward g)
152   (lambda (s/c)
153     (let rec (( $\lambda$  (g s/c)))
154       (cond ((null?  $\lambda$ ) '())
155             ((promise?  $\lambda$ ) (force  $\lambda$ ))
156             ((procedure?  $\lambda$ ) (lambda () (rec ( $\lambda$ ))))
157             (else (cons (car  $\lambda$ ) (rec (cdr  $\lambda$ )))))))
158
159 (define (bind  $\lambda$  g)
160   (cond
161     ((null?  $\lambda$ ) mzero)
162     ((procedure?  $\lambda$ ) (lambda () (bind ( $\lambda$ ) g)))
163     ((promise?  $\lambda$ ) (delay (bind (force  $\lambda$ ) (forward g))))
164     (else (mplus (g (car  $\lambda$ )) (bind (cdr  $\lambda$ ) g))))
165

```

The resulting 53 lines of code make up the functional core of our temporal  $\mu$ Kanren. To keep the subsequent code simple, we also introduce a few convenience functions for accessing delayed streams, here defined with the help of take-right and drop-right from SRFI-1.

```

169 (define (promised  $\lambda$ )
170   (take-right  $\lambda$  0))
171
172 (define (current  $\lambda$ )
173   (drop-right  $\lambda$  0))
174
175 (define (advance  $\lambda$ )
176   (let ((p (promised  $\lambda$ )))
177     (and p (force p))))
178

```

We are ready for a simple example that demonstrates its basic operation.

```

179 (define *db* 1)
180
181
182 (define r
183   ((call/fresh
184     (lambda (q)
185       (disj (== q *db*)
186             (next (== q *db*))))))
187   empty-state))
188 ;; => ((((#(0) . 1)) . 1) . #<promise>)
189
190 (set! *db* 2)
191
192 (advance r)
193 ;; => ((((#(0) . 2)) . 1))
194

```

The shunt-right mechanism turns out to have interesting properties for interleaving search even without referring to a stateful resource, for instance in ordering and grouping solutions, as the

following contrived example shows. Here and in what follows, will use the miniKanren wrappers from the original paper, likewise extended to handle delayed streams (the code is presented in Appendix A).

```
(define (inco x)
  (let r ((n 0))
    (disj (== x n) (next (r (+ n 1))))))

(define s
  (run* (q)
    (fresh (a b)
      (== q (list a b))
      (conj (inco a) (inco b)))))

(current s)
;; => ((0 0))

(current (advance s))
;; => ((0 1) (1 0) (1 1))

(current (advance (advance s)))
;; => ((1 2) (2 0) (2 1) (2 2) (0 2))

(current (advance (advance (advance s))))
;; => ((0 3) (2 3) (3 0) (3 1) (3 2) (3 3) (1 3))
```

#### 4 TOWARDS A TEMPORAL RELATIONAL PROGRAMMING

Linear temporal logic extends predicate logic with temporal modal operators referring to linear time. Two operators, Next (X) and Until (U), are generally defined as primitives and then used to construct a larger set of operators for conveniently reasoning about time. Several of these operators can be usefully recovered in temporal  $\mu$ Kanren using the next primitive. We will not develop a formally complete system here, but rather provide some heuristic examples to suggest how such a system might be developed.

Most temporal operators need to be defined recursively. When doing this, it is essential to control when the goal is constructed, since the goal might refer to a stateful resource. For temporal  $\mu$ Kanren we can do this by wrapping goals in a lambda expression which are evaluated at the appropriate time. The definition of precedes, modeled on Weak Until (W) and which stipulates that a goal g holds at least until a second goal h holds, shows how this is done.

```
(define (precedes* g* h*)
  (let ((g (g*)) (h (h*)))
    (disj h (conj g (next (precedes* g* h*)))))

(define-syntax precedes
  (syntax-rules ()
    ((_ g h) (precedes* (lambda () g) (lambda () h)))))
```

Translating another basic operator, Always (G), reveals one of the pitfalls of mapping temporal logic to relational programming. Our first impulse is to define it analogously to until.

```
(define (always* g*)
```

```

246   (let ((g (g*)))
247     (conj g (next (always* g*))))))
248
249 (define-syntax always
250   (syntax-rules ()
251     ((_ g) (always* (lambda () g)))))

```

252 A little experimentation, however, shows how this definition is problematic. A goal constructed  
 253 with `always` will return a promise so long as the goal it encloses holds, and the empty stream as  
 254 soon as it fails. While this behavior can be useful for a program that only wants to verify whether  
 255 a state still holds, `always` will never construct a mature stream of solutions, since clearly the goal  
 256 will never have been true *forever*, at least not in linear time.

257 There are two ways of getting around this problem. One is to modify the definition of `next` in  
 258 such a way as to allow for an `eot` (end-of-time) accessor that would essentially cut off the forward  
 259 recursion of delayed promises. Another approach is to define the weaker and impure `as-long-as`,  
 260 which constructs a stream of solutions to the goal `h` that continues as long as the goal `g` still holds.  
 261 This technique of writing ‘guarded’ goal constructors is reminiscent of the implementation of cuts  
 262 in `miniKanren`.

```

263 (define (as-long-as* g* h*)
264   (let ((g (g*)) (h (h*)))
265     (lambda (s/c)
266       (let (($ (g s/c)))
267         (if (null? $) mzero
268             (bind $ (disj h (next (as-long-as* g* h*))))))))))
269
270 (define-syntax as-long-as
271   (syntax-rules ()
272     ((_ g h) (as-long-as* (lambda () g) (lambda () h)))))

```

274 The Eventually (F) operator also has many useful applications, but presents a similar caveat re-  
 275 garding infinite time. Depending on the application, a guarded definition might be more appropriate,  
 276 but the following definition will still be useful in many contexts.

```

277 (define (eventually* g*)
278   (let ((g (g*)))
279     (disj g (next (eventually* g*))))))
280
281 (define-syntax eventually
282   (syntax-rules ()
283     ((_ g) (eventually* (lambda () g)))))
284
285 We can use it to define a strong Until.
286 (define-syntax until
287   (syntax-rules ()
288     ((_ g h) (conj (precedes g h) (eventually h)))))

```

289 Clearly this is only a beginning, and as the above examples suggest, what constitutes a *useful*  
 290 temporal relational programming language will be determined in part by the application domain.  
 291 Developing a complete version of such a system would be an obvious next step. Another direction  
 292 to pursue would be recovering temporal  $\mu$ Kanren’s functionality in a version of `miniKanren` with  
 293 constraints, namely `cKanren` [Alvis et al. 2011] or the extended  $\mu$ Kanren described in [Hemann  
 294

and Friedman 2015]. It would also be interesting to explore the integration of temporal logic with different representations of negation in miniKanren.

## 5 CALCULATING DELTAS WITH INCREMENTAL SEARCH

Now we turn to a practical application, and see how temporal  $\mu$ Kanren can be used to implement a simple data store with temporally-aware incremental search. We proceed by describing the implementation first, and then presenting the full code at the end of the section.

The motivation is as follows. In distributed systems such as a microservice architecture (for the specific context, see [Versteden and Pauwels 2016] and [Versteden and Pauwels 2018]) we often want to send push updates based on *deltas*, or entries that have been added to or removed from the database. In practice, this often means having a service that indiscriminately pushes all deltas to interested subscribers; it is then the responsibility of each subscriber to filter the deltas and determine which, if any, are relevant to its own operations.

If we can calculate deltas to specific queries, however, this whole process can be greatly refined. Here we describe how an RDF database (triple store) can be implemented using temporal  $\mu$ Kanren as its query language, that calculates deltas. By using the next constructor and a simple system of incremental indexes, we can store the final search positions for a query. Running a query will return both the current results and a delayed stream that when advanced will continue searching at the previous search-tree's leaves. Therefore, advancing the delayed stream after the database has been updated will return solutions that have been added to, or subtracted from the solution set.

An RDF database [Manola and Miller 2004] stores semantic facts, or triples,<sup>2</sup> made up of a *subject* (a URI), a *predicate* (a URI), and an *object* (a URI or a string, boolean or numeric literal).

```
<http://ex.org/people/1> <http://xmlns.com/foaf/0.1/name> "John"
```

A minimalist triple store can be implemented using three indexes, allowing for the retrieval of groups of triples matching a given query pattern: spo, pos, and osp, where s is the subject, p the predicate, and o the object. Here, however, we want to know *incremental indexes*: for a given s, we need the indexed p's in a form that can be easily compared to future p's to determine whether new keys have been added. We keep separate incremental indexes for each index level, storing the incrementals as a simple cons list along with a map for quick existence checking. For each index we use a hash array mapped trie<sup>3</sup> that allows for persistently storing successive states of the database through path copying.

Our database is therefore defined as a set of seven incremental indexes (null, s, sp, p, po, o, os) plus the spo index for full triples. Adding a triple to the database is a matter of consing each element to the appropriate incrementals lists and adding a truthy value (here #t, though this could also be a more meaningful identifier or timestamp) to the full spo index. When deleting a triple, we leave the indexes, but update the triple's value in spo to #f. Here are the incremental indexes after adding the triple <A> <B> <C>.

```
<∅ [ #f => (<A> ...) ] >
<s [ <A> => (<B> ...) ] >
<sp [ (<A> <B>) => (<C> ...) ] >
<p [ <B> => (<C> ...) ] >
<po [ (<B> <C>) => (<A> ...) ] >
<o [ <C> => (<A> ...) ] >
<os [ (<C> <A>) => (<B> ...) ] >
```

<sup>2</sup>Here we will only consider triples, though triple stores actually store quads, with the addition of the *graph* element.

<sup>3</sup>In the Chicken implementation presented here, we use Persistent Hash Maps [Heidkamp 2013], derived from Ian Price's Hash Array Mapped Tries [Price 2014].

```
344 <spo [ (<A> <B> <C>) => #t ] >
```

345 The main accessor is the goal constructor `triple-nolo` ('triple now or later') that descends  
 346 recursively through the incremental indexes one level at a time, constructing a stream with the  
 347 current indexes and saving the last search positions in a delayed stream. A dynamic parameter is  
 348 used to specify the current database state.

349 To see our data store in action, we consider the following query written in SPARQL 1.1 [Harris  
 350 and Seaborne 2013], the standard query language for RDF data stores. This query will be translated  
 351 into temporal  $\mu$ Kanren and run against successive states of the database.

```
352 SELECT ?o
353 WHERE {
354   <S> <P> ?o.
355   <Q> <R> ?o.
356 }
357
```

358 First, however, we define an empty database.

```
359 (define db0 (empty-db))
```

360 The SPARQL query is translated into the temporal  $\mu$ Kanren goal `r` using `triple-nolo`. To make  
 361 it clear what is going on, we keep track of the individual delta flags for each triple goal, though in  
 362 practice we usually want to know only if the solution was added (all `+`s) or removed (at least one `-`).  
 363 Since there are no solutions to our query, the program returns a delayed stream.

```
364 (define r
365   (parameterize ((latest-db db0))
366     (run* (q)
367       (fresh (o deltas d1 d2)
368         (== q `(<S> ,deltas ,o))
369         (== deltas `(<S> ,d1 ,d2))
370         (triple-nolo d1 '<S> '<P> o)
371         (triple-nolo d2 '<Q> '<R> o))))))
372 ;; => #<promise>
```

373 Now we can add some triples and advance the delayed stream.

```
374 (define db1
375   (add-triples db0 '((<S> <P> <01>)
376                     (<S> <P> <02>)
377                     (<Q> <R> <01>)
378                     (<A> <B> <C>))))
379
```

```
380 (parameterize ((latest-db db1))
381   (advance r))
382 ;; => (((+ +) <01>) . #<promise>)
```

384 If we delete a triple that contributed to one of our solutions, that solution will be returned with a  
 385 negative delta flag in the next iteration.

```
386 (define db2
387   (delete-triples db1 '((<S> <P> <01>))))
388
389 (parameterize ((latest-db db2))
390   (advance (advance r)))
391 ;; => (((+ -) <01>) . #<promise>)
```



Finally we can add that triple back along with some new solutions, to get some positive deltas.

```
(define db3
  (add-triples db2 '(<S> <P> <O1>)
                (<S> <P> <O3>)
                (<Q> <R> <O3>)
                (<S> <P> <M>)
                (<Q> <R> <M>))))

(parameterize ((latest-db db3))
  (advance (advance (advance r))))
;; => (((+ +) <O1>) ((+ +) <M>) ((+ +) <O3>) . #<promise>)
```

Here we have been proceeding linearly, but since the database states are persistent we can calculate deltas over any two states. Keeping track of a stream of states and indexing them on time will therefore give us a truly time-traveling data store.

Finally, here is the full code for the data store.

```
(define-syntax project
  (syntax-rules ()
    ((_ (x ...) g)
     (lambda (s/c)
       (let ((x (walk x (car s/c))) ...)
         (g s/c))))))

(define-record db s sp p po o os spo)

(define-record incrementals map list)

(define (empty-incrementals)
  (make-incrementals (persistent-map) '()))

(define (empty-db)
  (apply make-db
    (make-list 8 (persistent-map))))

(define latest-db (make-parameter (empty-db)))

(define (latest-incrementals accessor key)
  (incrementals-list
    (map-ref (accessor (latest-db)) key (empty-incrementals))))

(define (latest-triple s p o)
  (map-ref (db-spo (latest-db))
    (list s p o)))

(define (update-incrementals table key val)
  (let ((incrementals (map-ref table key (empty-incrementals))))
    (map-add table key
      (if (map-ref (incrementals-map incrementals) val)
```

```

442         incrementals
443         (make-incrementals
444         (map-add (incrementals-map incrementals) val #t)
445         (cons val (incrementals-list incrementals))))))
446
447 (define (update-triple table triple val)
448   (map-add table triple val))
449
450 (define (update-triples DB triples val)
451   (let loop ((triples triples)
452             (i/null (db-null DB))
453             (i/s (db-s DB))
454             (i/sp (db-sp DB))
455             (i/p (db-p DB))
456             (i/po (db-po DB))
457             (i/o (db-o DB))
458             (i/os (db-os DB))
459             (i/spo (db-spo DB)))
460     (if (null? triples)
461         (make-db i/null i/s i/sp i/p i/po i/o i/os i/spo)
462         (match (car triples)
463             ((s p o)
464              (loop (cdr triples)
                     (update-incrementals i/null #f s)
                     (update-incrementals i/s s p)
                     (update-incrementals i/sp (list s p) o)
                     (update-incrementals i/p p o)
                     (update-incrementals i/po (list p o) s)
                     (update-incrementals i/o o s)
                     (update-incrementals i/os (list o s) p)
                     (update-triple i/spo (list s p o) val))))))
463
464
465
466
467
468
469
470
471
472
473
474 (define (add-triples DB triples) (update-triples DB triples #t))
475
476 (define (delete-triples DB triples) (update-triples DB triples #f))
477
478 (define (add-triple s p o)
479   (add-triples `((,s ,p ,o))))
480
481 (define (delete-triple s p o)
482   (delete-triples `((,s ,p ,o))))
483
484 (define (triple-nolo delta s p o)
485   (let ((mkstrm (lambda (var accessor key)
486                  (let* ((get-incrementals (lambda ()
487                                              (latest-incrementals accessor key)))
488                        (initial-indexes (get-incrementals)))
489                    (let stream ((indexes initial-indexes)

```

```

491         (ref '()) (next-ref initial-indexes))
492     (if (equal? indexes ref)
493         (next
494         (let ((vals (get-incrementals)))
495             (stream vals next-ref vals)))
496         (disj
497         (conj (= var (car indexes))
498             (project (delta s p o) (triple-nolo delta s p o)))
499         (stream (cdr indexes) ref next-ref))))))
500 (cond ((and (var? s) (var? p) (var? o)) (mkstrm s db-null #f))
501       ((and (var? s) (var? p))          (mkstrm s db-o o))
502       ((and (var? s) (var? o))          (mkstrm o db-p p))
503       ((and (var? p) (var? o))          (mkstrm p db-s s))
504       ((var? s)                         (mkstrm s db-po (list p o)))
505       ((var? p)                         (mkstrm p db-os (list o s)))
506       ((var? o)                         (mkstrm o db-sp (list s p)))
507       (else
508        (let leaf ((ref #f))
509            (let ((v (latest-triple s p o)))
510                (cond ((eq? v ref) (next (leaf v)))
511                      (v          (disj (= delta '+) (next (leaf v))))
512                      (else       (disj (= delta '-') (next (leaf v))))))))))

```

Though this reduced example is clearly far from a production-ready system, we hope we have demonstrated the one of the practical uses of temporal relational programming. Indeed, we intend to pursue these tools in an industrial setting, and the full implementation will be a laboratory for other practical applications of miniKanren, such as drawing on [Byrd et al. 2012] to compile SPARQL to miniKanren and using search ordering as explored in [Swords and Friedman 2013] to aid with query optimization.

## A APPENDIX: MINIKANREN WRAPPERS

Here we adapt the miniKanren control operators described in [Hemann and Friedman 2013] for delayed streams. Unlike the original paper, we do not use `Zzz` to wrap goals in `conj+` and `disj+` since this makes it difficult of control the goal construction time at different levels of nested delays, a problem when referring to stateful resources. This difficulty should be addressed in future implementations of temporal  $\mu$ Kanren.

```

527 (define-syntax Zzz
528   (syntax-rules ()
529     ((_ g) (lambda (s/c) (lambda () (g s/c))))))
530
531 (define-syntax conj+
532   (syntax-rules ()
533     ((_ g) g)
534     ((_ g0 g ...) (conj g0 (conj+ g ...))))))
535
536 (define-syntax disj+
537   (syntax-rules ()
538     ((_ g) g)

```

```

540      (( _ g0 g ...) (disj g0 (disj+ g ...))))))
541
542 (define-syntax fresh
543   (syntax-rules ()
544     (( _ () g0 g ...) (conj+ g0 g ...))
545     (( _ (x0 x ...) g0 g ...)
546       (call/fresh
547         (lambda (x0)
548           (fresh (x ...) g0 g ...))))))
549
550 (define-syntax conde
551   (syntax-rules ()
552     (( _ (g0 g ...) ...) (disj+ (conj+ g0 g ...) ...))))
553
554 (define-syntax run
555   (syntax-rules ()
556     (( _ n (x ...) g0 g ...)
557       (let r ((k n) ($ (take n (call/goal (fresh (x ...) g0 g ...))))))
558         (cond ((null? $) '())
559               ((promise? $) (delay (r (- k 1) (take k (force $)))))
560               (else (cons (reify-1st (car $))
561                           (r (- k 1) (cdr $)))))))
562
563 (define-syntax run*
564   (syntax-rules ()
565     (( _ (x ...) g0 g ...)
566       (let r (($ (take-all (call/goal (fresh (x ...) g0 g ...))))))
567         (cond ((null? $) '())
568               ((promise? $) (delay (r (take-all (force $)))))
569               (else (cons (reify-1st (car $))
570                           (r (cdr $)))))))
571
572 (define empty-state '(() . 0))
573
574 (define (call/goal g) (g empty-state))
575
576 (define (pull $)
577   (cond ((procedure? $) (pull ($)))
578         ((promise? $) $)
579         (else $)))
580
581 (define (take-all $)
582   (let (($ (pull $)))
583     (cond ((null? $) '())
584           ((promise? $) $)
585           (else (cons (car $) (take-all (cdr $))))))
586
587 (define (take n $)
588

```

```

589     (if (zero? n) '()
590         (let (($ (pull $)))
591             (cond ((null? $) '())
592                   ((promise? $) $)
593                   (else (cons (car $) (take (- n 1) (cdr $)))))))
594
595 (define (reify-1st s/c)
596   (let ((v (walk* (var 0) (car s/c))))
597     (walk* v (reify-s v '()))))
598
599 (define (walk* v s)
600   (let ((v (walk v s)))
601     (cond
602       ((var? v) v)
603       ((pair? v) (cons (walk* (car v) s)
604                         (walk* (cdr v) s)))
605       (else v))))
606
607 (define (reify-s v s)
608   (let ((v (walk v s)))
609     (cond
610       ((var? v)
611        (let ((n (reify-name (length s))))
612          (cons `(. ,n) s)))
613       ((pair? v) (reify-s (cdr v) (reify-s (car v) s)))
614       (else s))))
615
616 (define (reify-name n)
617   (string->symbol
618     (string-append "_" "." (number->string n))))
619
620 (define (fresh/nf n f)
621   (letrec
622     ((app-f/v*
623      (lambda (n v*)
624        (cond
625          ((zero? n) (apply f (reverse v*)))
626          (else (call/fresh
627                  (lambda (x)
628                    (app-f/v* (- n 1) (cons x v*))))))))
629     (app-f/v* n '()))))

```

## REFERENCES

- Claire E. Alvis, Jeremiah J. Willcock, Kyle M. Carter, William E. Byrd, and Daniel P. Friedman. 2011. cKanren: miniKanren with Constraints. *Proceedings of the 2011 Workshop on Scheme and Functional Programming*.
- William E. Byrd. 2009. *Relational programming in miniKanren: techniques, applications, and implementations*. Ph.D. Dissertation. Bloomington, IN, USA.
- William E. Byrd, Eric Holk, and Daniel P. Friedman. 2012. miniKanren, Live and Untagged: Quine Generation via Relational Interpreters (Programming Pearl). In *Proceedings of the 2012 Annual Workshop on Scheme and Functional Programming*

- (Scheme '12). ACM, New York, NY, USA, 8–29. <https://doi.org/10.1145/2661103.2661105>
- Daniel P. Friedman, William E. Byrd, and Oleg Kiselyov. 2005. *The Reasoned Schemer*. The MIT Press, Cambridge, MA, USA.
- Steve Harris and Andy Seaborne. 2013. *SPARQL 1.1 Query Language*. Technical Report. <https://www.w3.org/TR/sparql11-query/>
- Moritz Heidkamp. 2013. persistent-hash-map. (2013). <http://wiki.call-cc.org/eggref/4/persistent-hash-map>
- Jason Hemann and Daniel P. Friedman. 2013.  $\mu$ Kanren: A Minimal Functional Core for Relational Programming. *Proceedings of the 2013 Workshop on Scheme and Functional Programming*.
- Jason Hemann and Daniel P. Friedman. 2014. microKanren. (2014). <https://github.com/jasonhemann/microKanren>
- Jason Hemann and Daniel P. Friedman. 2015. A Framework for Extending microKanren with Constraints. *Proceedings of the 2015 Workshop on Scheme and Functional Programming*.
- Frank Manola and Eric Miller. 2004. *RDF Primer*. Technical Report. <https://www.w3.org/TR/2004/REC-rdf-primer-20040210/>
- Mehmet A. Orgun and Wanli Ma. 1994. An Overview of Temporal and Modal Logic Programming. In *Proc. First Int. Conf. on Temporal Logic - LNAI 827*. Springer-Verlag, 445–479.
- Ian A. Price. 2014. Purely Functional Data Structures in Scheme. (2014). <https://github.com/ijp/pfds/blob/master/hamts.sls>
- Cameron Swords and Daniel Friedman. 2013. rKanren: Guided search in miniKanren. In *Proceedings of the 2013 Workshop on Scheme and Functional Programming*.
- Aad Versteden and Erika Pauwels. 2016. State-of-the-art Web Applications using Microservices and Linked Data. In *Proceedings of the 4th Workshop on Services and Applications over Linked APIs and Data*.
- Aad Versteden and Erika Pauwels. 2018. mu.semte.ch: State-of-the-art Web Applications fuelled by Linked Data. (2018). Presentation at the Extended Semantic Web Conference 2018, Heraklion, Greece.